

Exploring Diabetic Health Indicators Dataset

GROUP 6

TEAM MEMBERS:

JHARANA ADHIKARI: C0927442

UMA MAHESHWARI : C0928443

YASASWIN PALUKURI: C0928450



Overview



Why Diabetes Health Indicators Dataset?



Characteristics of Diabetic Health Indicator



Lifecycle of Big Data Analytics



Utilization of Analysis Result



Conclusion



Reference



Why Diabetes Health Indicators Dataset?



- Research on diabetes is essential to enhancing patient outcomes and lowering complications.
- It lessens the strain on healthcare systems by assisting in the development of cost-effective strategies.
- Research improves diabetics' overall health, early detection, and treatment.

Why Diabetes Health Indicators Dataset?

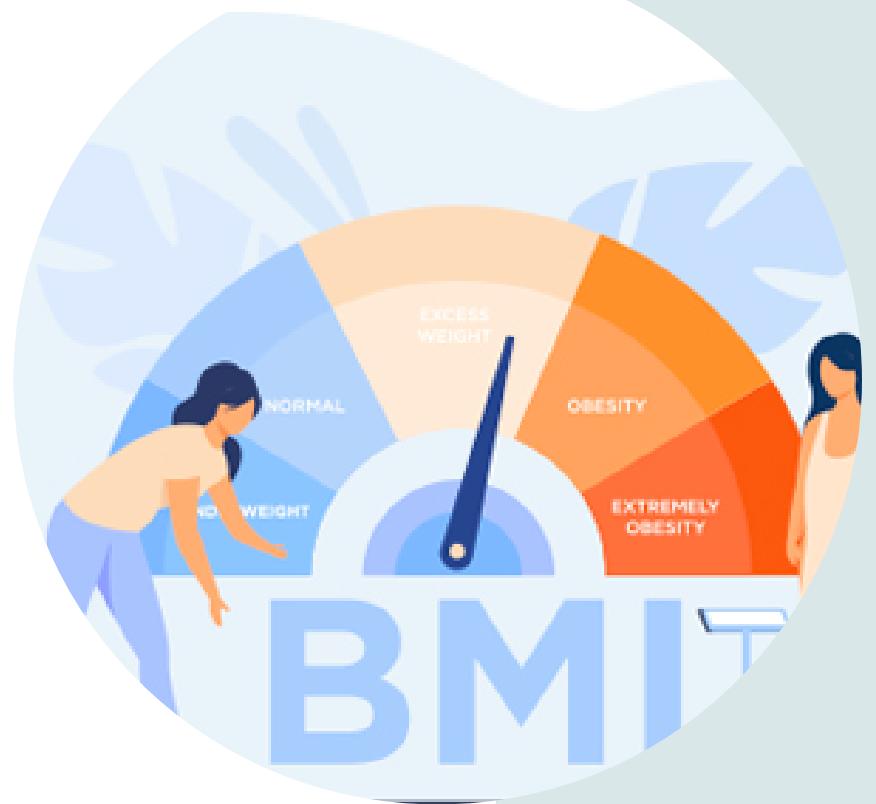
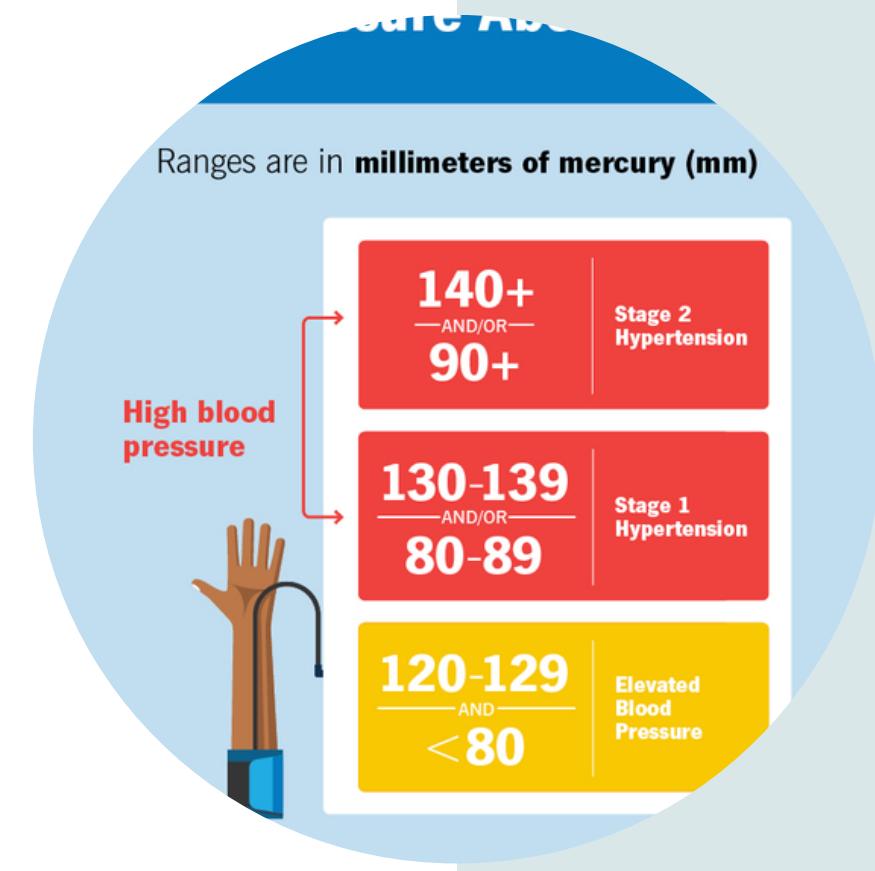


- Rich dataset with diverse health indicators provides ample features for building robust predictive models.
- Machine learning algorithms can leverage the complexity of the dataset to identify patterns and relationships, aiding in diabetes risk prediction and prevention strategies.

Characteristics of Diabetic Health

Indicator Datasets:

- **Diabetes_012:** Indicates 0 for no diabetes/pregnancy, 1 for prediabetes, and 2 for diabetes; crucial variable for analysis.
- **HighBP:** High blood pressure indicator; relevant for diabetes risk assessment.
- **HighChol:** High cholesterol indicator; associated with diabetes complications.
- **CholCheck:** Cholesterol check frequency; reflects health awareness.
- **BMI:** Body Mass Index; correlates with diabetes risk.



Characteristics of Diabetic Health

Indicator Datasets:

- **Smoker:** Smoking status; impacts diabetes susceptibility.
- **Stroke:** Stroke history; potential complication of diabetes.
- **HeartDiseaseorAttack:** Heart disease history; linked to diabetes.
- **PhysActivity:** Physical activity level; influences diabetes prevention.
- **Fruits:** Fruit consumption; dietary factor in diabetes management.
- **Veggies:** Vegetable consumption; dietary factor affecting overall health and potentially diabetes risk.



Characteristics of Diabetic Health Indicator Datasets:

- **HvyAlcoholConsump:** Heavy alcohol consumption; lifestyle factor potentially impacting diabetes risk and overall health.
 - **AnyHealthcare:** Healthcare utilization; important for diabetes management.
 - **NoDocbcCost:** Doctor visit affordability; affects healthcare access.
 - **GenHlth:** General health perception; subjective health indicator.
 - **MentHlth:** Mental health status; psychological aspect of diabetes.





Characteristics of Diabetic Health

Indicator Datasets:

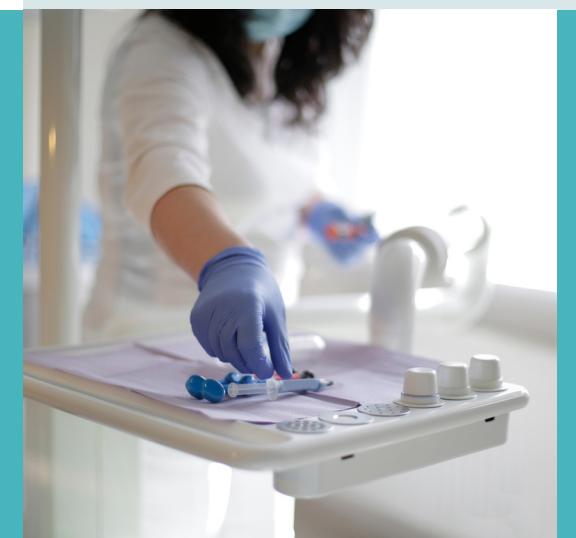
- 
- **PhysHlth: Physical health status; relevant for diabetes complications.**
 - **DiffWalk: Difficulty walking; potential diabetes symptom.**
 - **Sex: Gender; demographic variable in diabetes research.**
 - **Age: Age of respondents; correlates with diabetes prevalence.**
 - **Education: Educational attainment; socio-economic factor in diabetes.**
 - **Income: Household income; socio-economic determinant of health.**

BUSINESS CASE EVALUATION

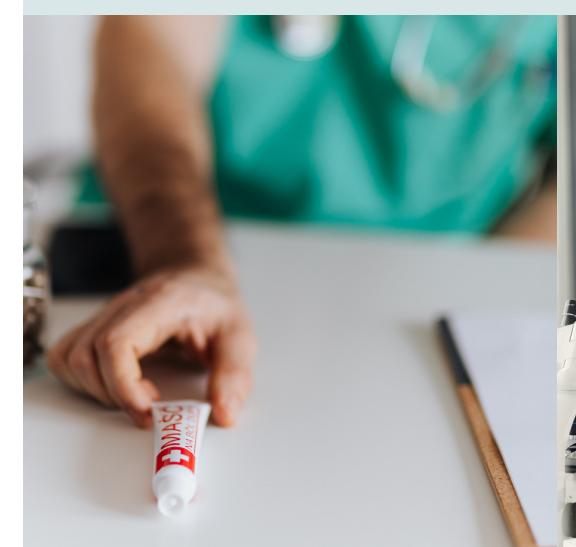
- CDC reports **34.2 million Americans with diabetes and 88 million with prediabetes as of 2018.**
- Shockingly, **1 in 5 diabetics and 8 in 10 predabetics are unaware of their risk, highlighting the urgency for early detection.**
- Type II diabetes, most common, costs roughly **\$327 billion** for diagnosed cases, and nearly **\$400 billion** annually including undiagnosed cases, disproportionately affecting lower socioeconomic groups.



Public
Health
Impact



Diabetes
Risk
Detection



DATA IDENTIFICATION

- **Dataset Overview:** Clean dataset comprising 253,680 survey responses to the CDC's BRFSS2015.
- **Target Variable:** Diabetes_012 with 3 classes: 0 for no diabetes/pregnancy, 1 for prediabetes, and 2 for diabetes
- **Feature Variables:** 21 feature variables available for analysis, providing insights into diabetes risk factors and indicators.



DATA ACQUISITION AND FILTERING

<https://www.kaggle.com/datasets/alexteboul/diabetes-health-indicators-dataset/data>



Diabetes Health Indicators Dataset

253,680 survey responses from cleaned BRFSS 2015 + balanced dataset

 kaggle.com

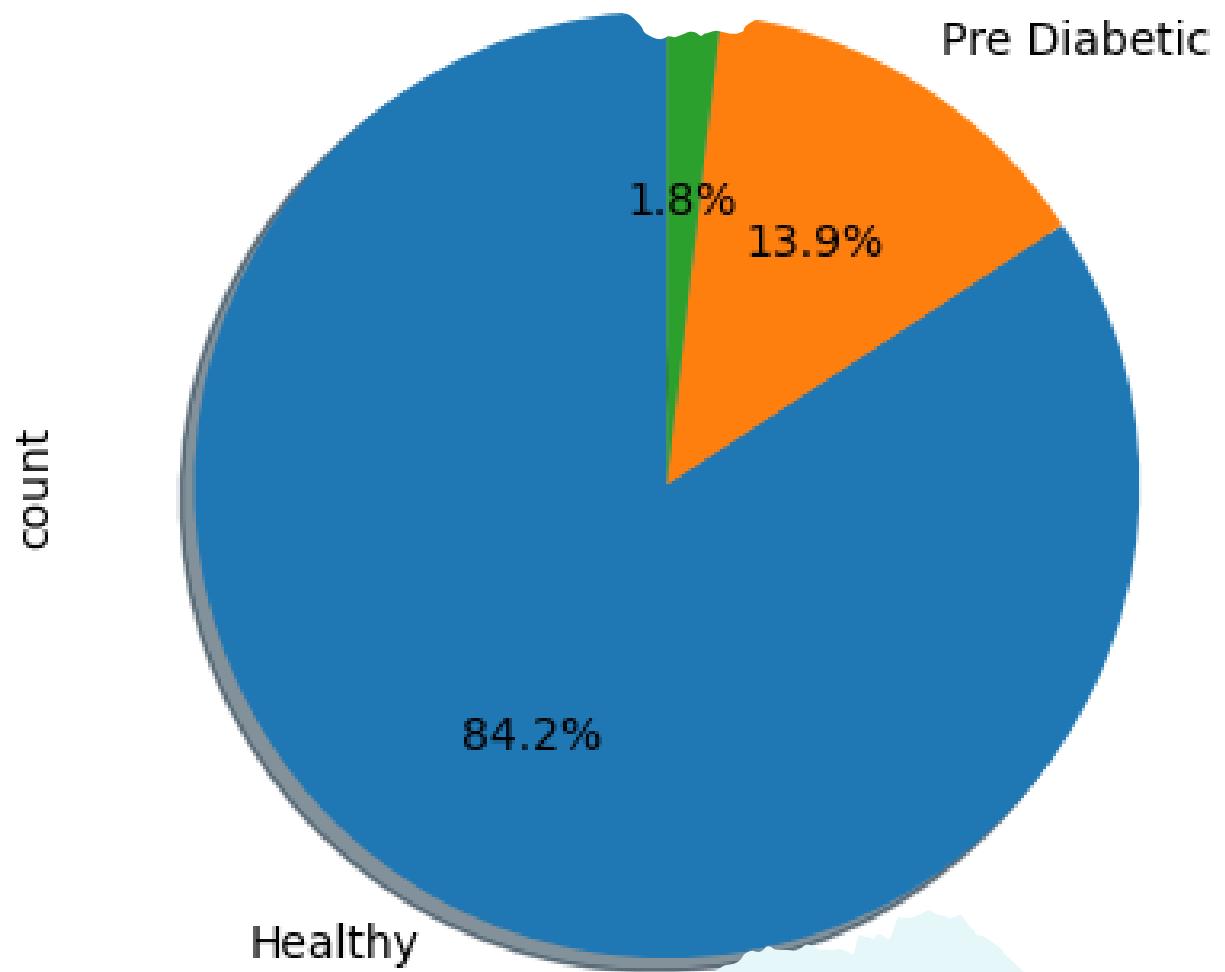
DATA EXTRACTION

- **Important characteristics include lifestyle, medical history, demographics, and symptoms.**
- **Make use of the pandas library in Python to extract specific data.**
- **Use info, head, and shape to comprehend datasets.**

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 253680 entries, 0 to 253679
Data columns (total 22 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Diabetes_012    253680 non-null  float64 
 1   HighBP          253680 non-null  float64 
 2   HighChol        253680 non-null  float64 
 3   CholCheck       253680 non-null  float64 
 4   BMI             253680 non-null  float64 
 5   Smoker          253680 non-null  float64 
 6   Stroke          253680 non-null  float64 
 7   HeartDiseaseorAttack  253680 non-null  float64 
 8   PhysActivity    253680 non-null  float64 
 9   Fruits          253680 non-null  float64 
 10  Veggies         253680 non-null  float64 
 11  HvyAlcoholConsump  253680 non-null  float64 
 12  AnyHealthcare   253680 non-null  float64 
 13  NoDocbcCost    253680 non-null  float64 
 14  GenHlth         253680 non-null  float64 
 15  MentHlth        253680 non-null  float64 
 16  PhysHlth        253680 non-null  float64 
 17  DiffWalk        253680 non-null  float64 
 18  Sex              253680 non-null  float64 
 19  Age              253680 non-null  float64 
 20  Education        253680 non-null  float64 
 21  Income           253680 non-null  float64 
dtypes: float64(22)
memory usage: 42.6 MB
```

DATA AGGREGATION & REPRESENTATION



MIN MAX SCALER

```
Columns
```

```
df2.head()
```

```
[]:
```

| | Diabetes_012 | HighBP | HighChol | BMI | Smoker | Stroke | HeartDiseaseorAttack | PhysActivity | Fruits | Veggies | HvyAlcoholConsump | GenHlth | MentHlth | P |
|---|--------------|--------|----------|------|--------|--------|----------------------|--------------|--------|---------|-------------------|---------|----------|---|
| 0 | 0.0 | 1.0 | 1.0 | 40.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 5.0 | 18.0 | |
| 1 | 0.0 | 0.0 | 0.0 | 25.0 | 1.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 3.0 | 0.0 | |
| 2 | 0.0 | 1.0 | 1.0 | 28.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 5.0 | 30.0 | |
| 3 | 0.0 | 1.0 | 0.0 | 27.0 | 0.0 | 0.0 | 0.0 | 1.0 | 1.0 | 1.0 | 0.0 | 2.0 | 0.0 | |
| 4 | 0.0 | 1.0 | 1.0 | 24.0 | 0.0 | 0.0 | 0.0 | 1.0 | 1.0 | 1.0 | 0.0 | 2.0 | 3.0 | |

```
# copy Dataframe columns
```

```
numerical_columns = ["Diabetes_012", "BMI", "GenHlth", "MentHlth", "PhysHlth", "Age"]
```

```
df3 = df2.copy(deep=True)
```

```
# Selecting numerical columns (excluding binary/boolean col
```

```
# Initialize the MinMaxScaler
```

```
scaler = MinMaxScaler()
```

```
# Fit and transform the numerical features
```

```
df3[numerical_columns] = scaler.fit_transform(df3[numerical_columns])
```

DATA ANALYSIS

- EDA (Exploratory Data Analysis)
- Finding the correlation between variables
- Splitting the data into test and train subsets
- Build a machine-learning model
- Perform a cross-validation technique



PREDIABETES in the United States in 2018

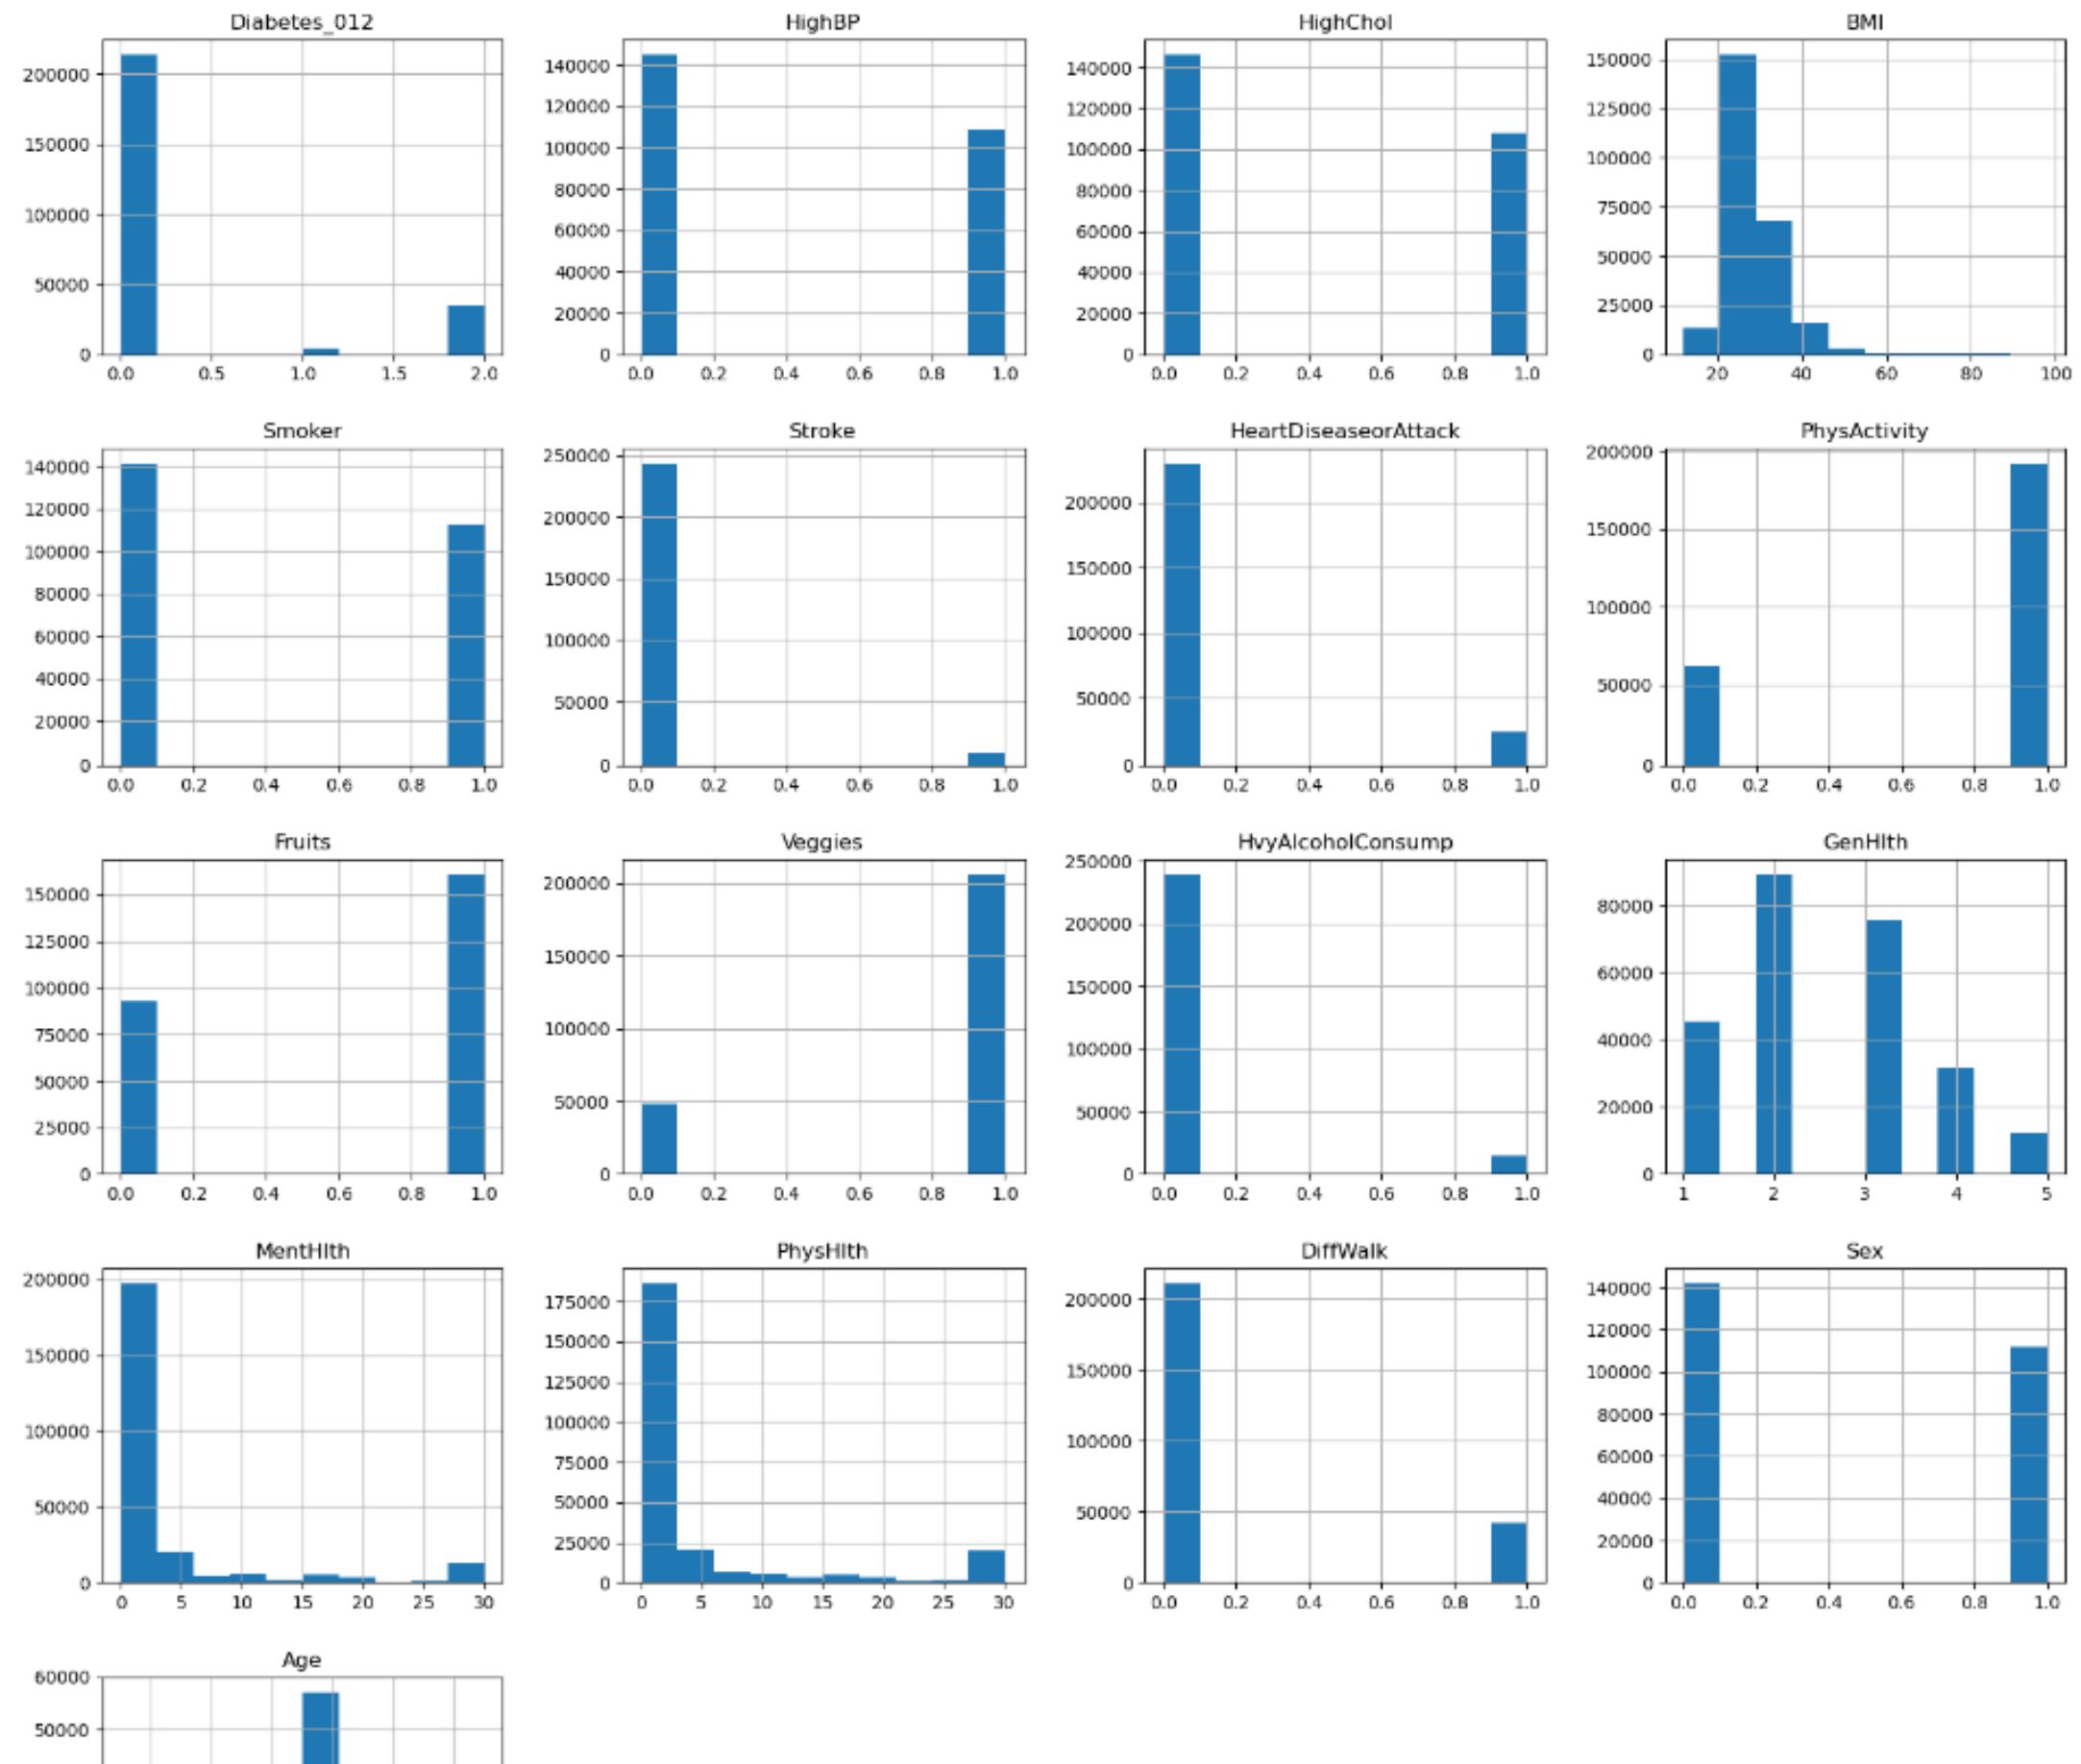
88 million
adults in the U.S.
—more than **1 in 3**—
were living with prediabetes
in 2018

84%
did not know they had it

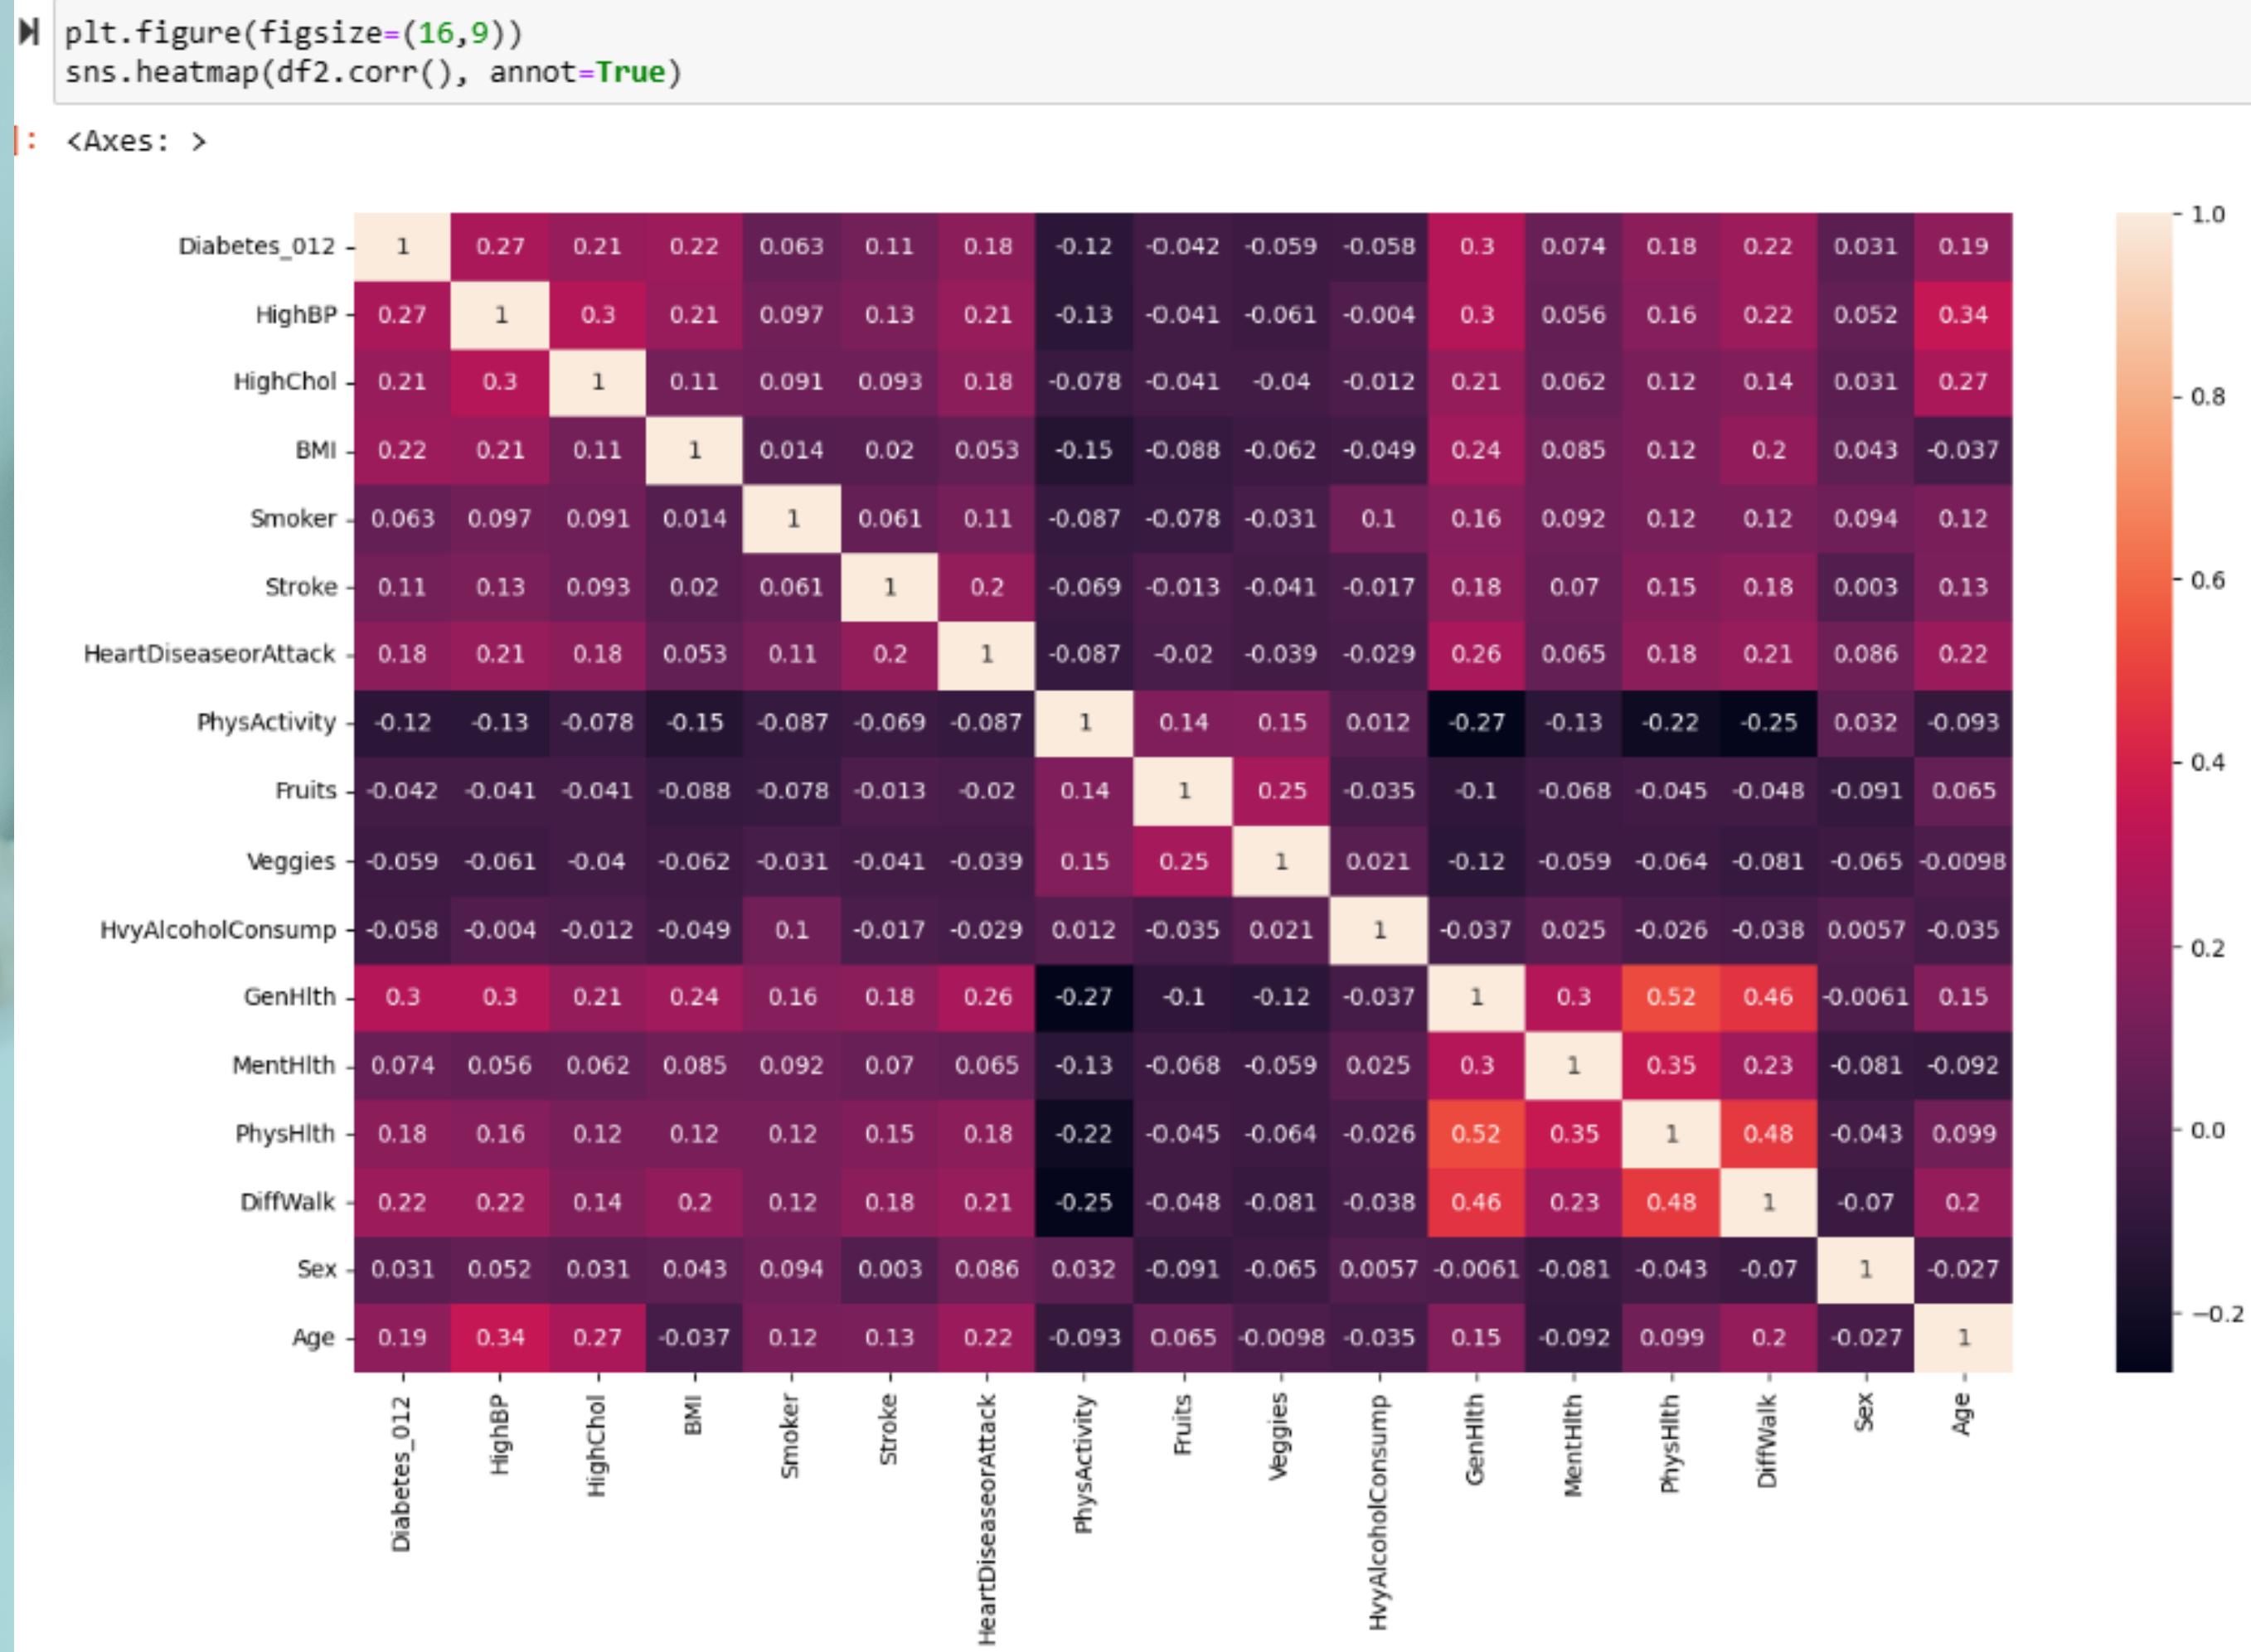
EDA (EXPLORATORY DATA ANALYSIS)

7.1 EDA (Exploratory Data Analysis)

```
df2.hist(figsize = (20,20))  
plt.show()
```



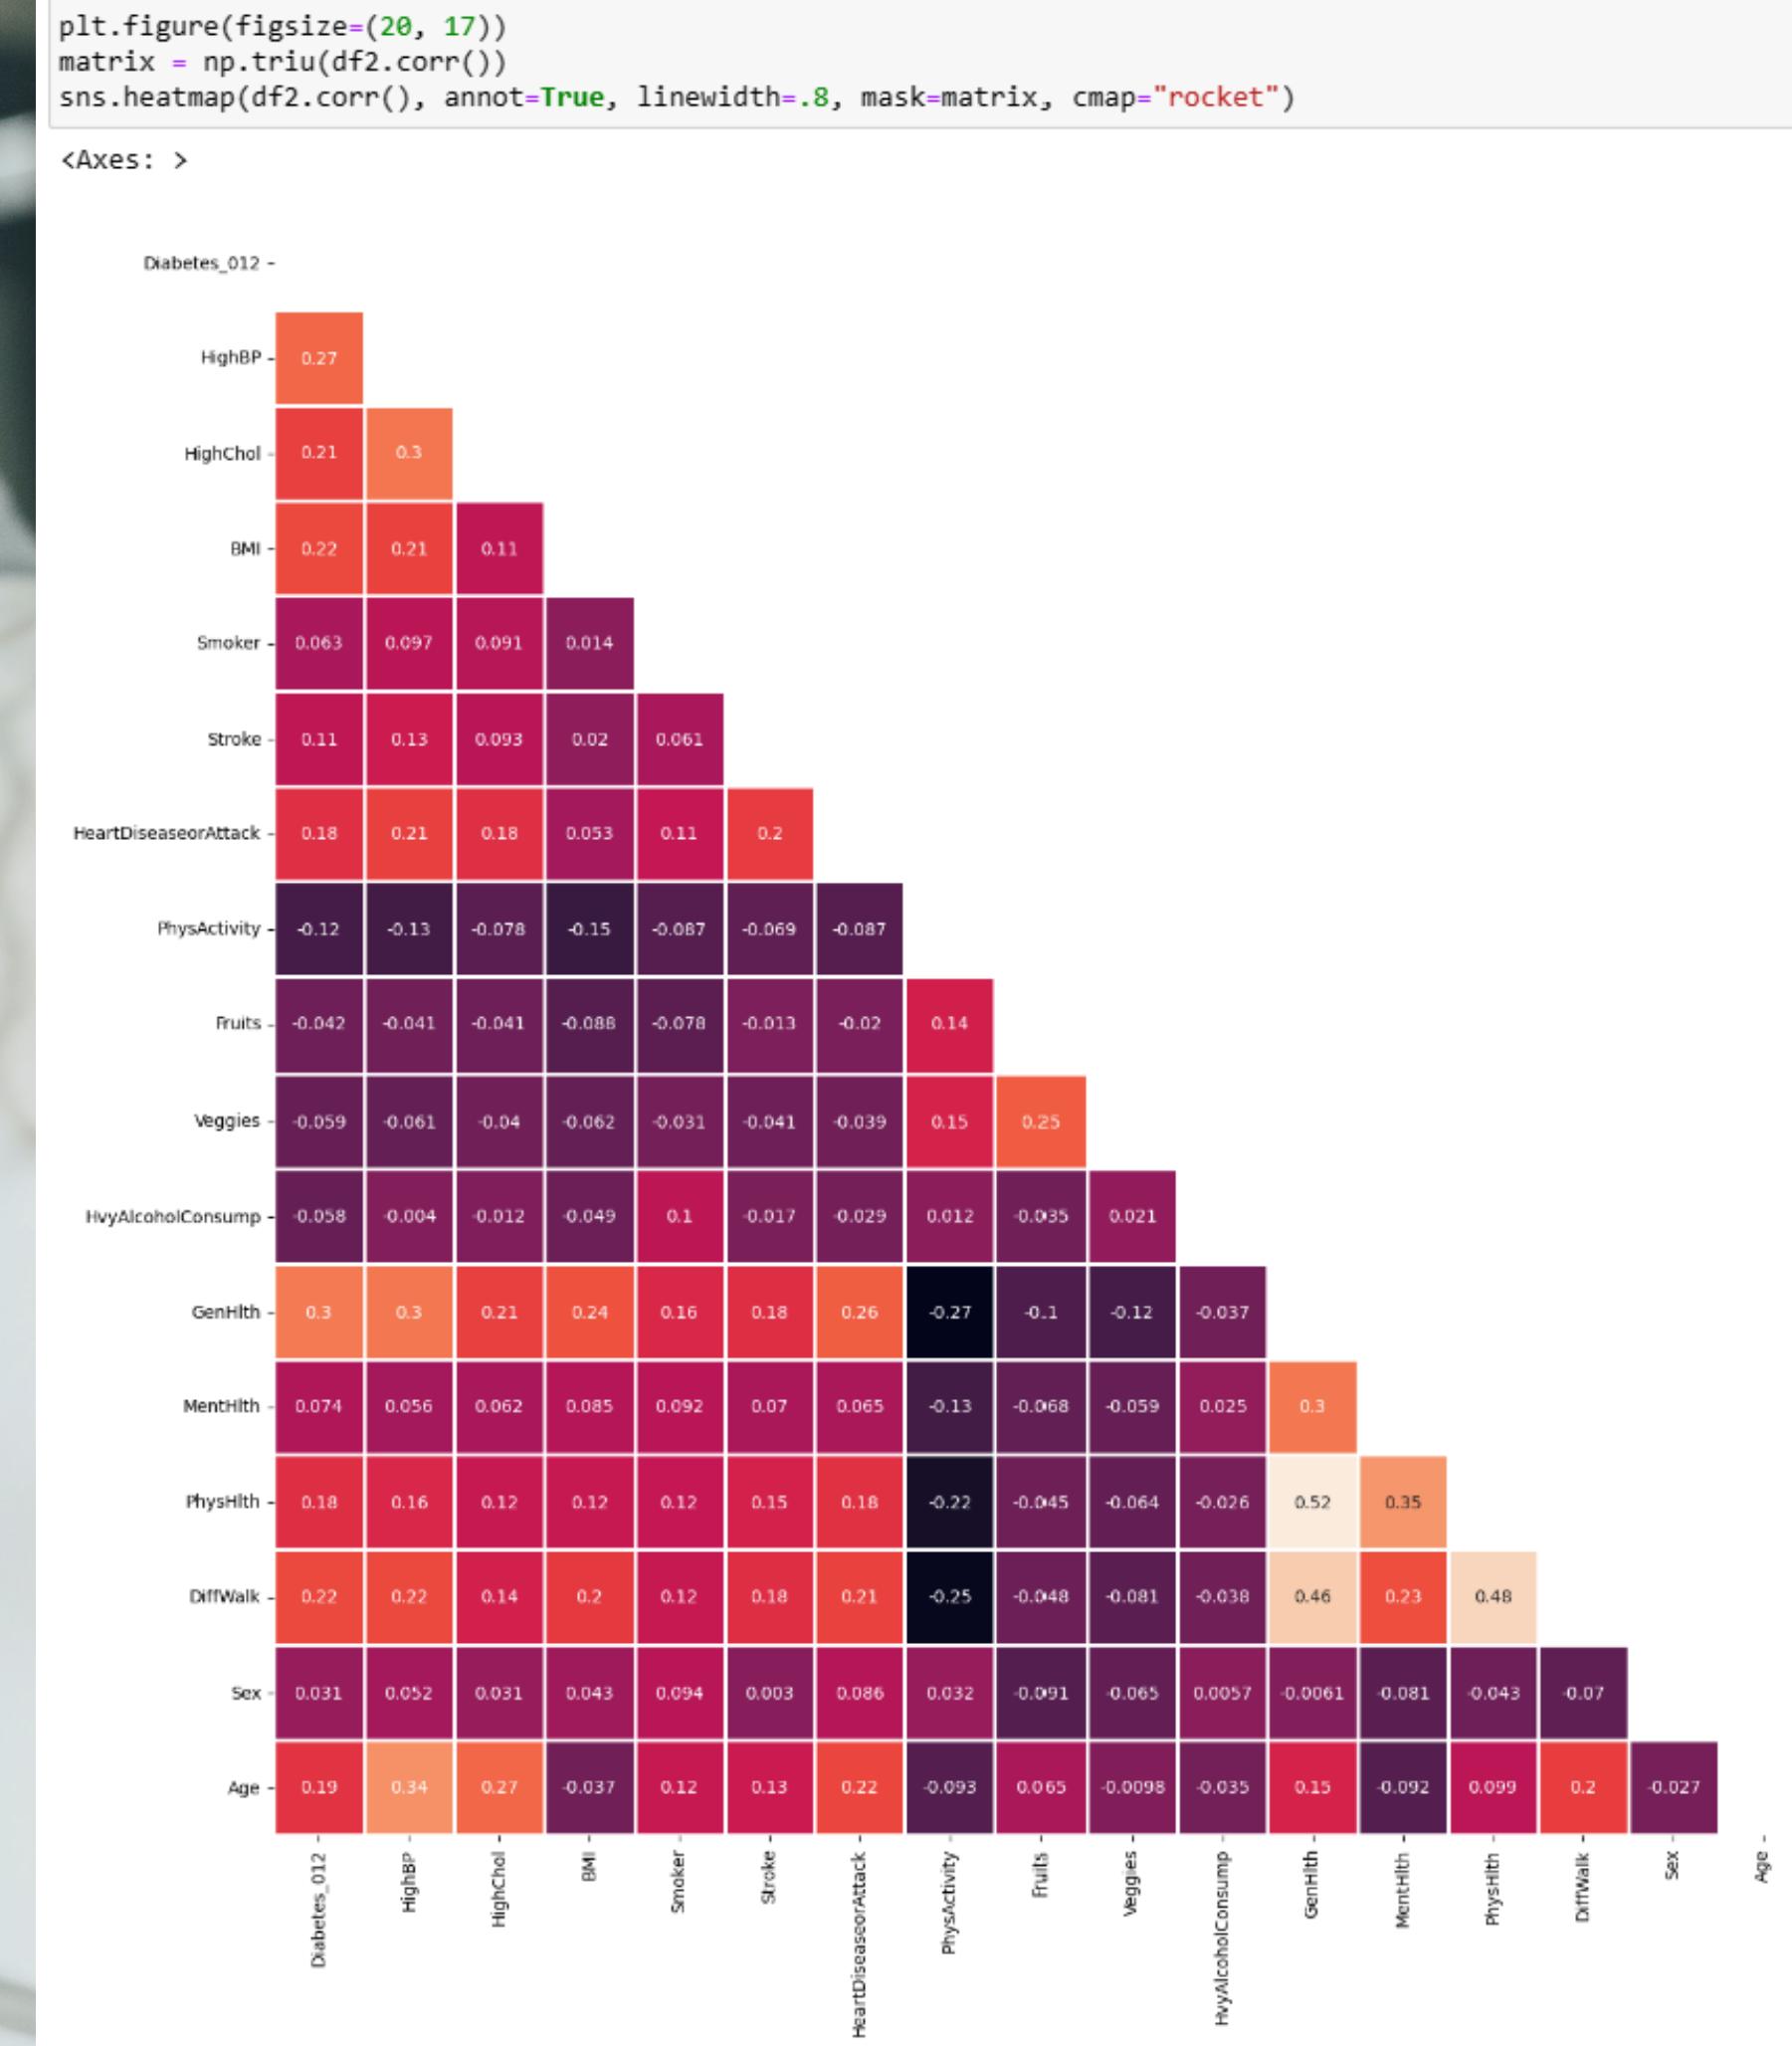
Finding the correlation between variables



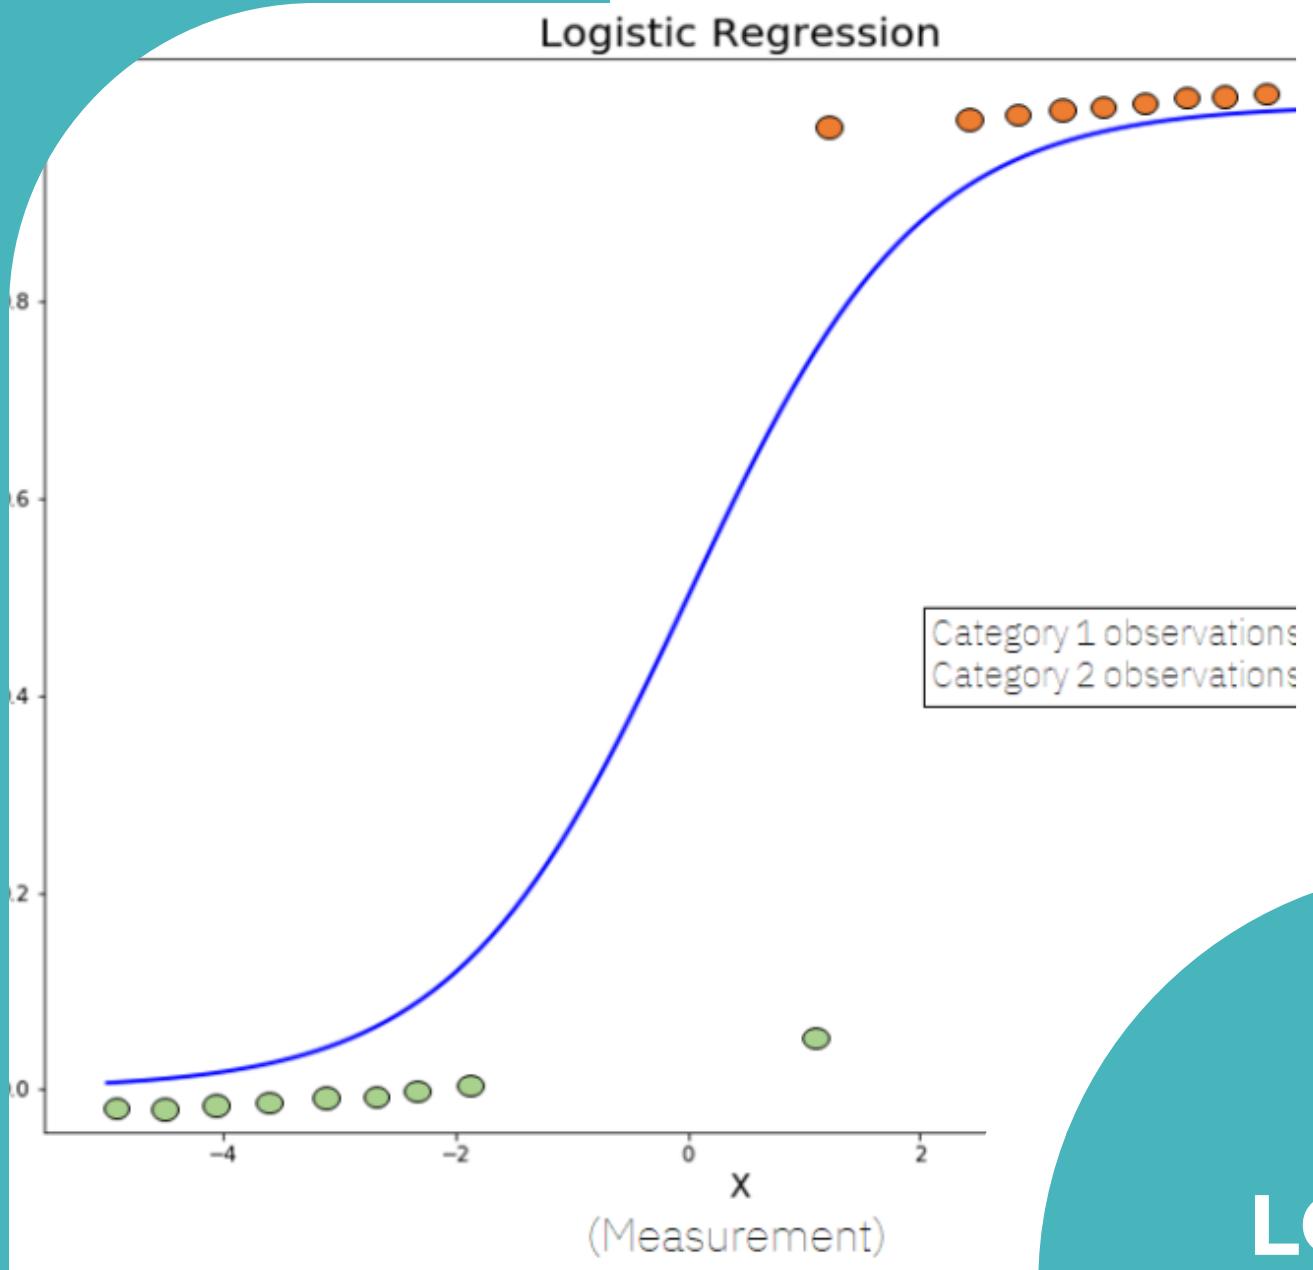
Finding the correlation between variables

```
hig_corr = df2.corr()  
hig_corr_features = hig_corr.index[abs(hig_corr["Diabetes_012"]) >= 0.2]  
hig_corr_features
```

```
Index(['Diabetes_012', 'HighBP', 'HighChol', 'BMI', 'GenHlth', 'DiffWalk'],
```



LOGISTIC REGRESSION



```
#Creating a Logistic regression classifier
log_model = linear_model.LogisticRegression()

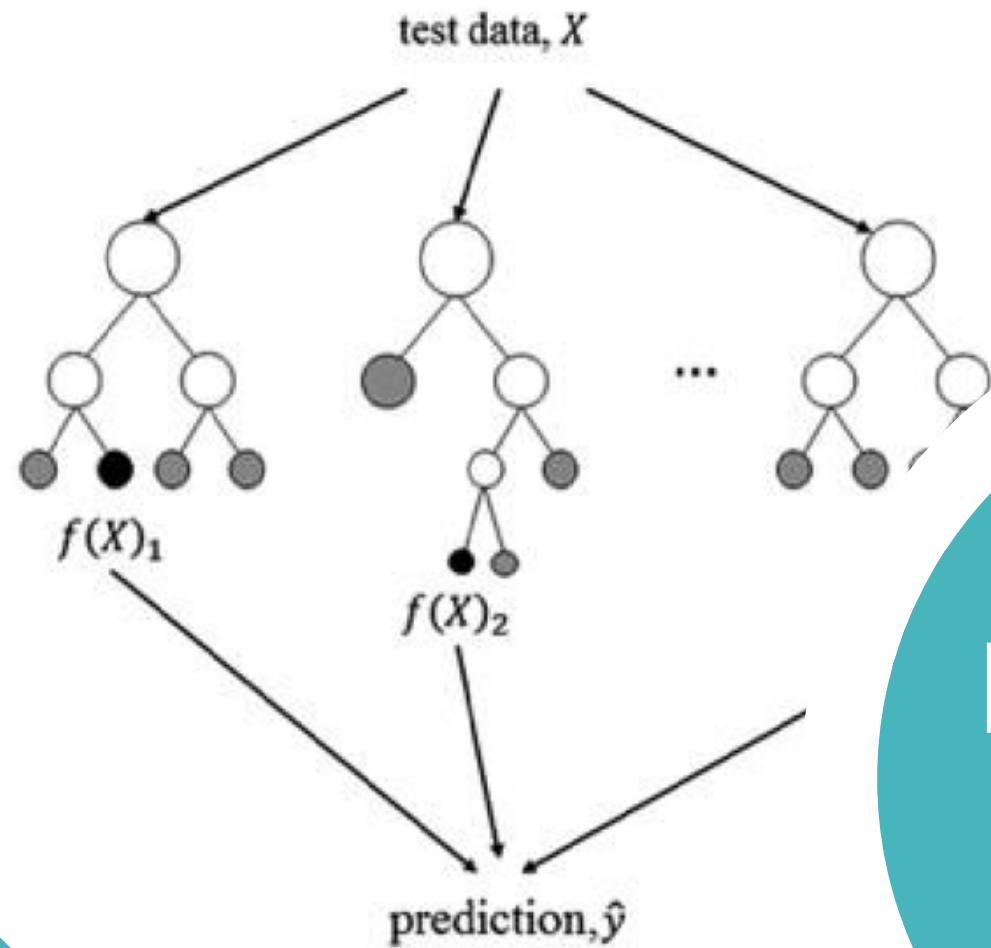
# Fitting Random Forest to the Training set
log_model.fit(X_train, y_train)

# Predicting the Test set results
y_pred = log_model.predict(X_test)

# Evaluate the model
accuracy= accuracy_score(y_test, y_pred)
print(f"Accuracy: {round(accuracy*100,2)} %", "\n")
```

Accuracy: 84.86 %

Testing



RANDOM FOREST

```
# Creating a Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_
                                         _state=42)

# Fitting Random Forest to the Training set
rf_classifier.fit(X_rf_train, y_rf_train)

# Predicting the Test set results
y_pred_rf = rf_classifier.predict(X_rf_test)

# Evaluate the model
accuracy_rf = accuracy_score(y_rf_test, y_pred_rf)

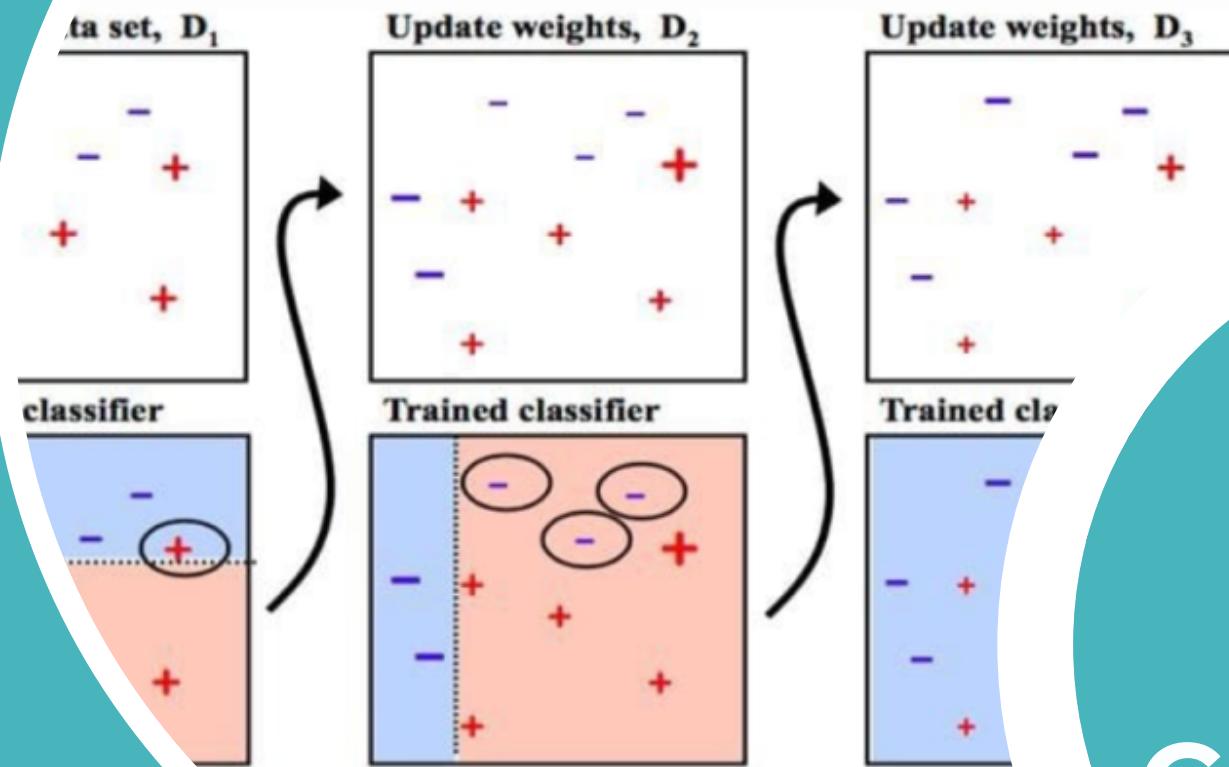
print("\nRandom Forest prediction results:", "\n")
print(f"Accuracy: {round(accuracy_rf*100,2)} %", "\n")
```

Random Forest prediction results

Accuracy: 83.48 %

Accuracy: 83.48 %

What is AdaBoost



ADA
BOOST
CLASSIFIER

```
ada_model = AdaBoostClassifier(n_estimators=50, random_state=42)  
  
# Fitting Ada Boost to the Training set  
ada_model.fit(X_train, y_train)  
  
# Predicting the Test set results  
y_pred_adaboost = ada_model.predict(X_test)  
  
# Evaluate the model  
accuracy = accuracy_score(y_test, y_pred_adaboost)  
print(f"Accuracy: {round(accuracy*100,2)} %", "\n")
```

Accuracy: 85.11 %

Accuracy: 85.11 %

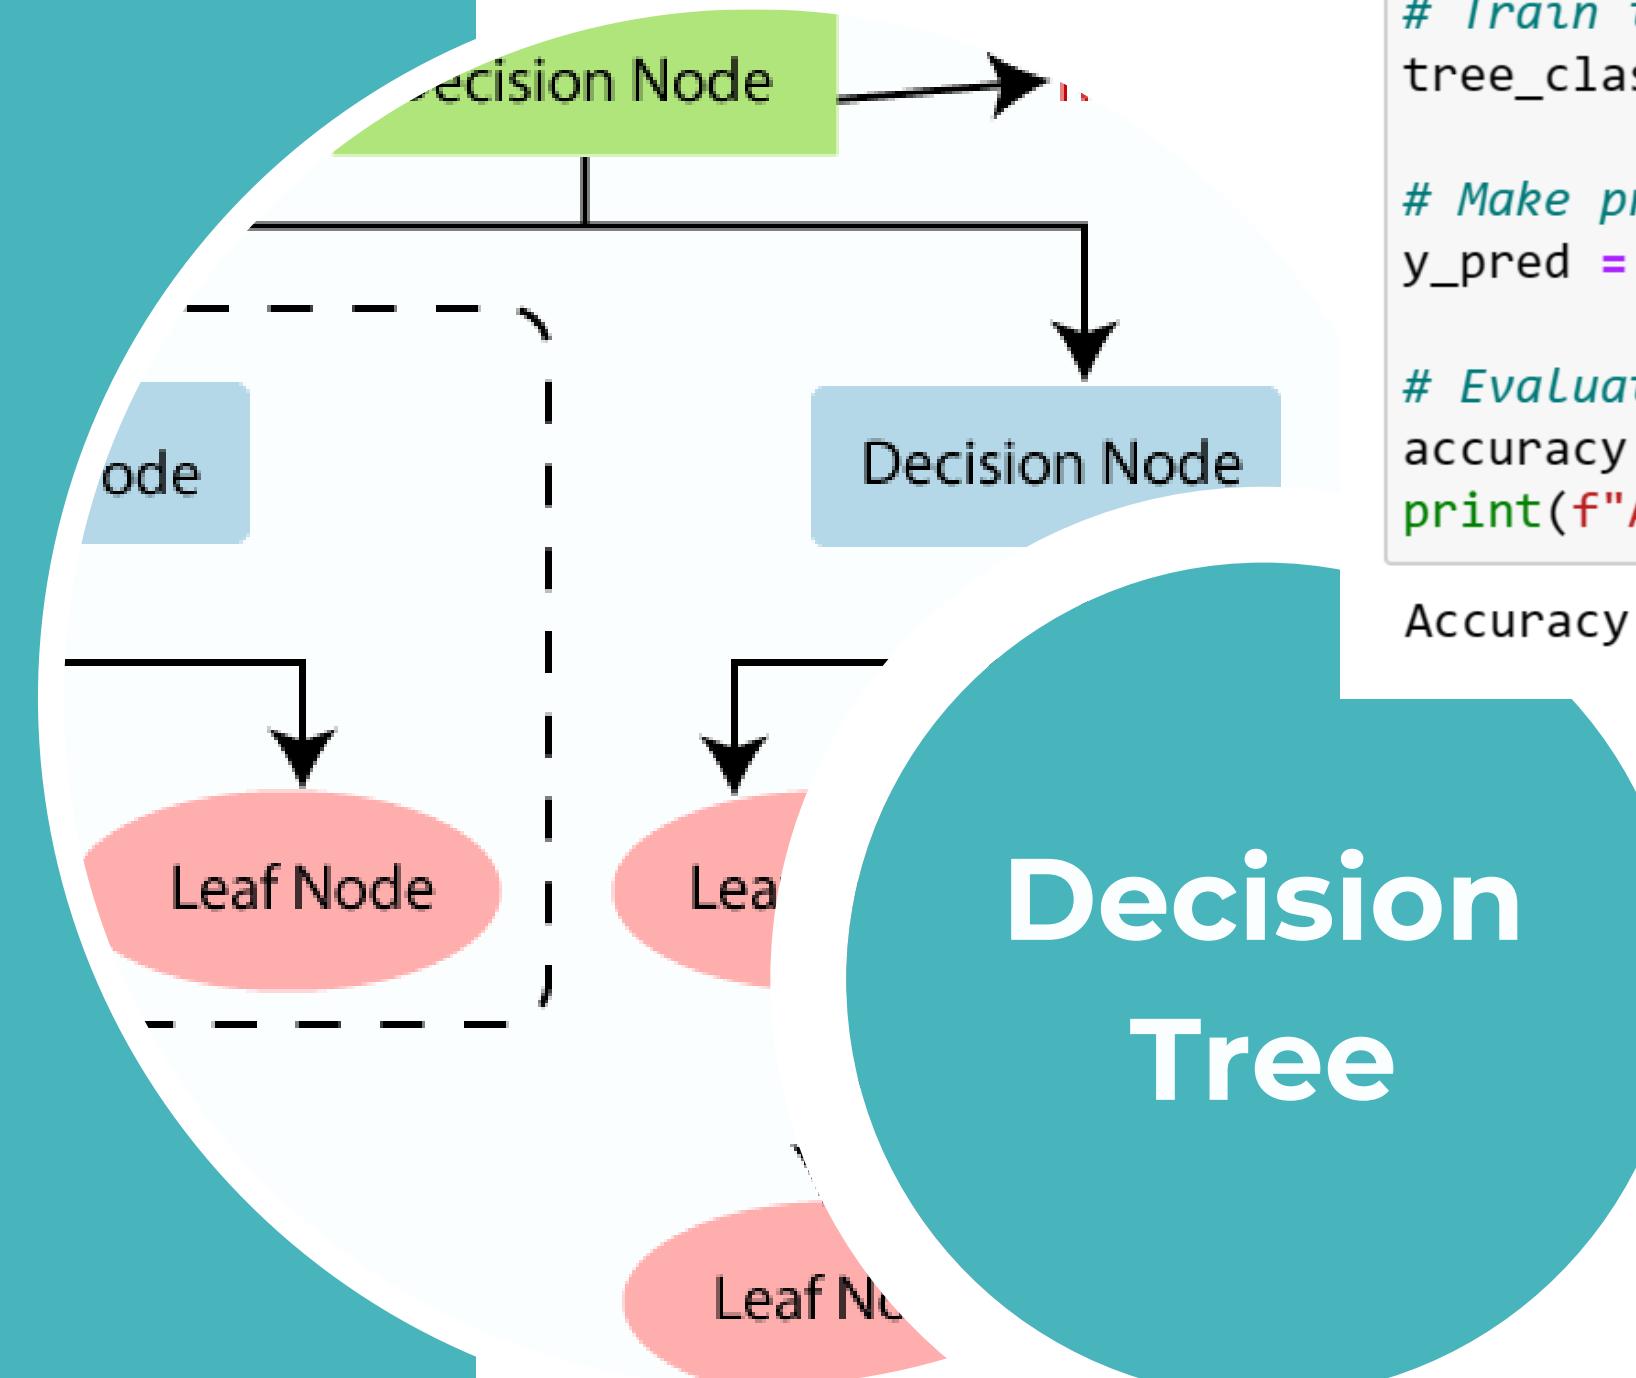
GRADIENT BOOSTING



```
| gbc_model = GradientBoostingClassifier(n_estimators=100,  
  
# Fit the GBC model to the training data  
gbc_model.fit(X_train, y_train)  
  
# Predicting the Test set results  
y_pred_gbm = gbc_model.predict(X_test)  
  
# Evaluate the model  
accuracy = accuracy_score(y_test, y_pred_gbm)  
print(f"Accuracy: {round(accuracy*100,2)} %", "\n")
```

Accuracy: 85.12 %

E. Decision Tree



```
# Create a DecisionTreeClassifier object
tree_classifier = DecisionTreeClassifier()

# Train the Decision Tree classifier
tree_classifier.fit(X_train, y_train)

# Make predictions on the testing data
y_pred = tree_classifier.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy*100:.2f}%")
```

Accuracy: 79.25%

Accuracy: 79.25%

Semi-supervised learning

```
X = df.drop('Diabetes_012', axis=1)
y = df['Diabetes_012']

# Just for demonstration, let's convert it to a semi-supervised setup by 'hiding' some
# Here, we arbitrarily mark 75% of the dataset as 'unlabeled' (-1)
mask = np.random.rand(len(y)) < 0.75
y[mask] = -1

# Splitting the dataset into labeled and unlabeled sets
X_labeled = X[y != -1]
y_labeled = y[y != -1]
X_unlabeled = X[y == -1]

# Confirm the distribution of classes
print(y_labeled.value_counts())

# Scaling the features (assuming X is not yet scaled for this subset)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Reducing the dataset size
X_small = X_scaled[:10000] # Adjust the number based on available memory
y_small = y[:10000].copy() # Make a copy to avoid warnings when changing this array

# Simulate semi-supervised setup by 'hiding' some
mask = np.random.rand(len(y_small)) < 0.75
y_small[mask] = -1 # Marking 75% of the data
```

```
Diabetes_012
0.0    53611
2.0     8905
1.0    1141
Name: count, dtype: int64
```

Label Propagation

```
# Splitting the dataset
X_train, X_test, y_train, y_test = train_test_split(X_small, y_small, test_size=0.2, random_state=42)

# Apply semi-supervised Learning: Label Propagation
# Initialize Label Propagation with a fixed random state
label_prop_model = LabelPropagation()

label_prop_model.fit(X_train, y_train)

# Predict on the test set
y_pred = label_prop_model.predict(X_test)

# Since we've used -1 for unlabeled, evaluate only on originally labeled data in y_test
mask_test = y_test != -1
print("Accuracy on test set:", accuracy_score(y_test[mask_test], y_pred[mask_test]))
print("Classification report on test set:\n", classification_report(y_test[mask_test], y_pred[mask_test]))
```

Accuracy on test set: 0.7818181818181819

Classification report on test set:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.85 | 0.89 | 0.87 | 90 |
| 1.0 | 0.00 | 0.00 | 0.00 | 2 |
| 2.0 | 0.43 | 0.33 | 0.38 | 18 |
| accuracy | | | 0.78 | 110 |
| macro avg | 0.43 | 0.41 | 0.41 | 110 |
| weighted avg | 0.77 | 0.78 | 0.77 | 110 |

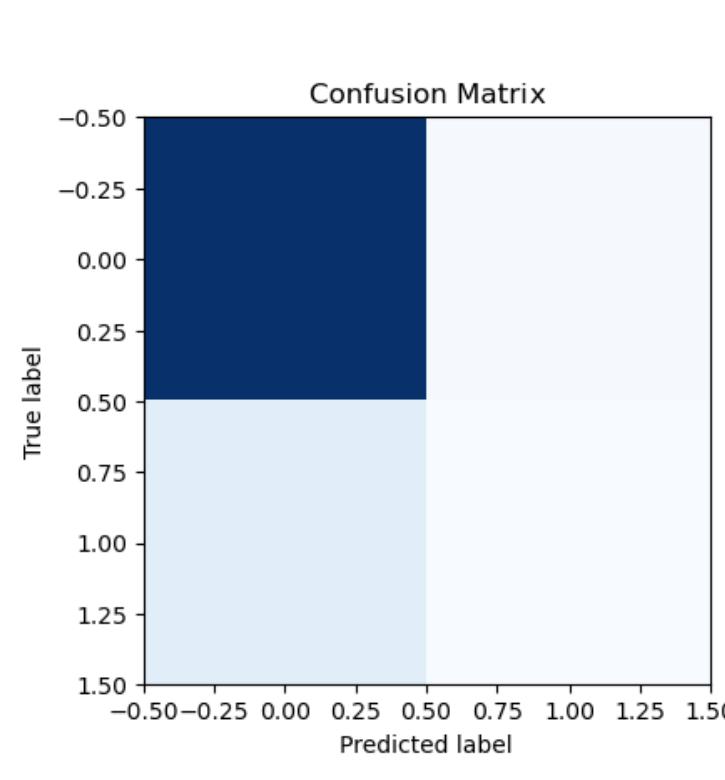
Data Visualization

Confusion matrix

```
#data visualization
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, classification_report, confusion_matrix
from sklearn.metrics import precision_recall_curve

# Plot confusion matrix
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
cm = confusion_matrix(y_rf_test, rf_classifier.predict(X_rf_test))
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.colorbar()
plt.title('Confusion Matrix')
plt.ylabel('True label')
plt.xlabel('Predicted label')

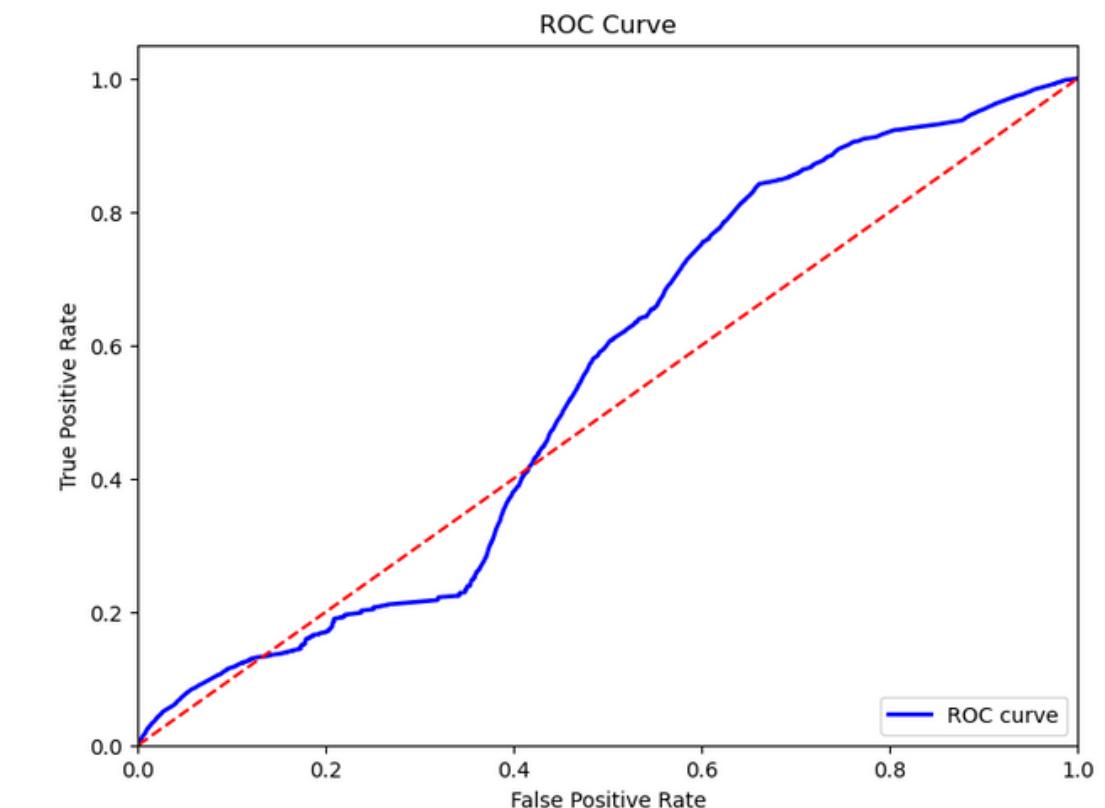
28]: Text(0.5, 0, 'Predicted label')
```



Roc Curve

```
# Compute ROC curve
fpr, tpr, thresholds = roc_curve(y_test, gbc_model.predict_proba(X_rf_test)[:, 1])

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve')
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc="lower right")
plt.show()
```

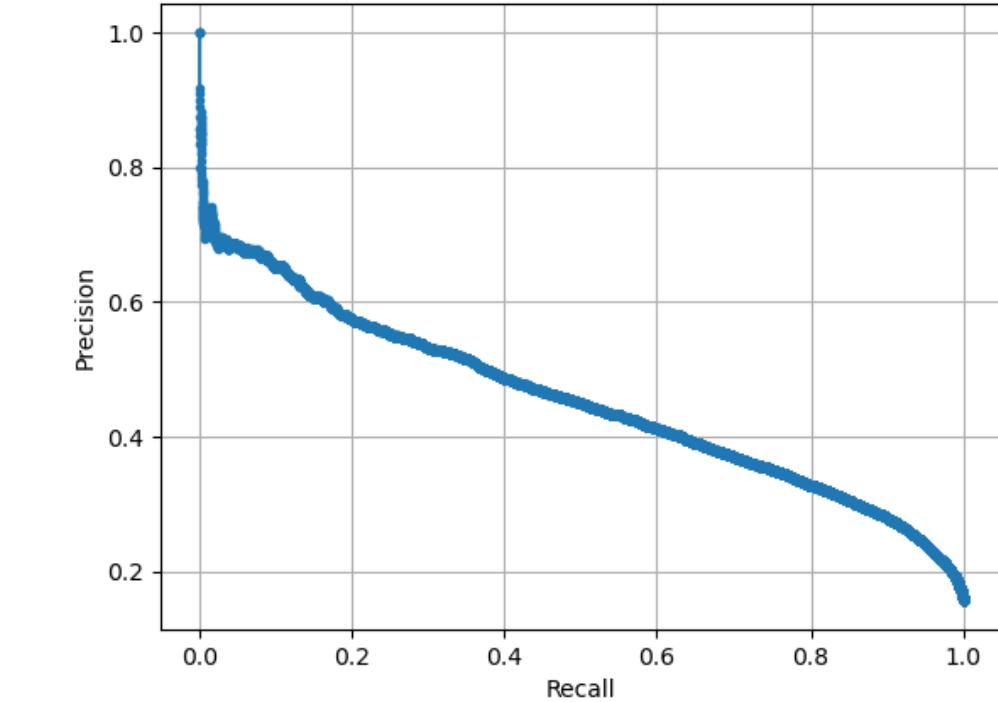


Precision-Recall Curve

```
# Compute precision-recall pairs for different probability thresholds
precisions, recalls, thresholds = precision_recall_curve(y_test, gbc_model.predict_proba(X_rf_test)[:, 1])

# Plot precision-recall curve
plt.plot(recalls, precisions, marker='.')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.grid(True)
plt.show()
```

Precision-Recall Curve



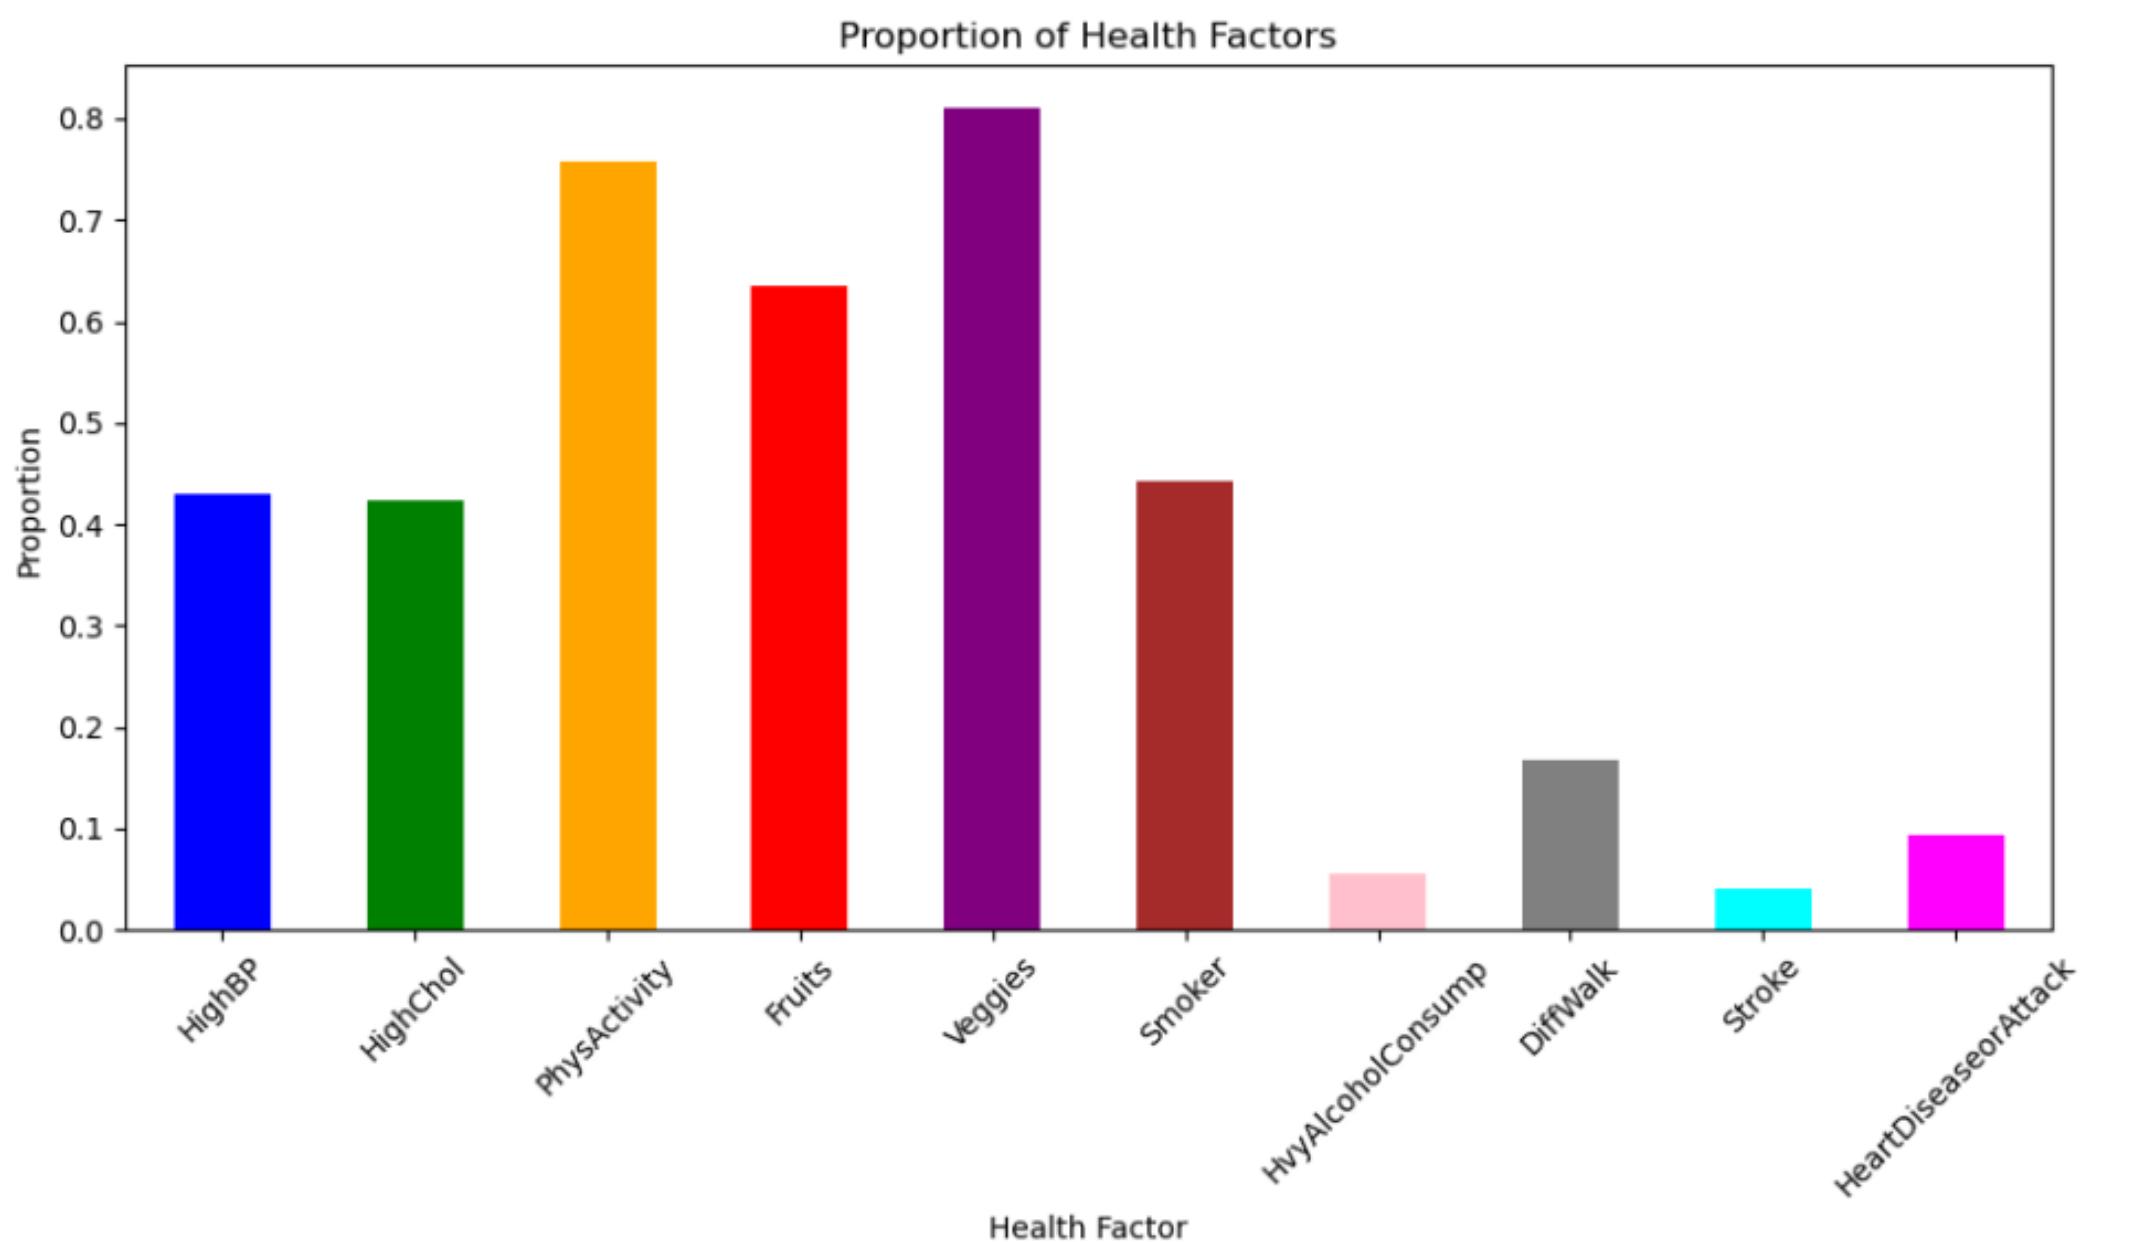
```

combinedBarChart = [
    'HighBP', 'HighChol', 'PhysActivity', 'Fruits', 'Veggies',
    'Smoker', 'HvyAlcoholConsump', 'DiffWalk', 'Stroke', 'HeartDiseaseorAttack'
]
df4 = df[combinedBarChart]

health_factors = df4.sum() / len(df4)

plt.figure(figsize=(10, 6))
health_factors.plot(kind='bar', color=['blue', 'green', 'orange', 'red', 'purple', 'brown', 'pink', 'gray', 'cyan', 'magenta'])
plt.title('Proportion of Health Factors')
plt.xlabel('Health Factor')
plt.ylabel('Proportion')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



Data Visualization



Diabetes
Health
Indicators

Utilization of Analysis Result

- **Identify high-risk individuals for diabetes:**
 - Consider demographic factors and health-related features
 - High BP, high cholesterol, high BMI, low general health increase risk
- **Allocate healthcare resources efficiently:**
 - Intensive interventions for high-risk individuals
 - Standard care for low-risk individuals
- **Target interventions towards:**
 - Weight management
 - Blood pressure control
 - Smoking cessation



Conclusion

- The datasets with key figures of diabetes health indicators allow to spot high-risk individuals and innovate in the way available resources are used in healthcare management.
- By making use of these datasets the decisions are made in the right way and thus prevention and management of diabetes are enhanced, and public health outcomes become more of a benefit to the community as a whole.



Diabetes
Health
Indicators

REFERENCES



Diabetes Health Indicators Dataset

253,680 survey responses from cleaned BRFSS 2015
+ balanced dataset

[kaggle.com](https://www.kaggle.com)



Diabetes - Symptoms and causes

Learn more about the different types of this blood sugar disorder, who's at risk and how each type can be treated.

 Mayo Clinic

Q&A



*Thank
You*

